

# Data Science Methods for Clean Energy Research

---

Week 9 L1: Decision Trees

~~Feb 26, 2017~~ → March 1 2017

UNIVERSITY *of* WASHINGTON



# Outline

---

- > Quick review from last time
- > Intro tree methods
  - Definition and properties
  - Regression trees
  - Classifier trees
- > Python project
- > Methods to improve decision trees
  - Bagging, boosting, random forests
- > Python example
- > Wrap up



# Topics last time

---

- > K-fold cross validation
- > Subset selection
- > Shrinkage/regularization methods
  - Ridge regression
  - LASSO regression

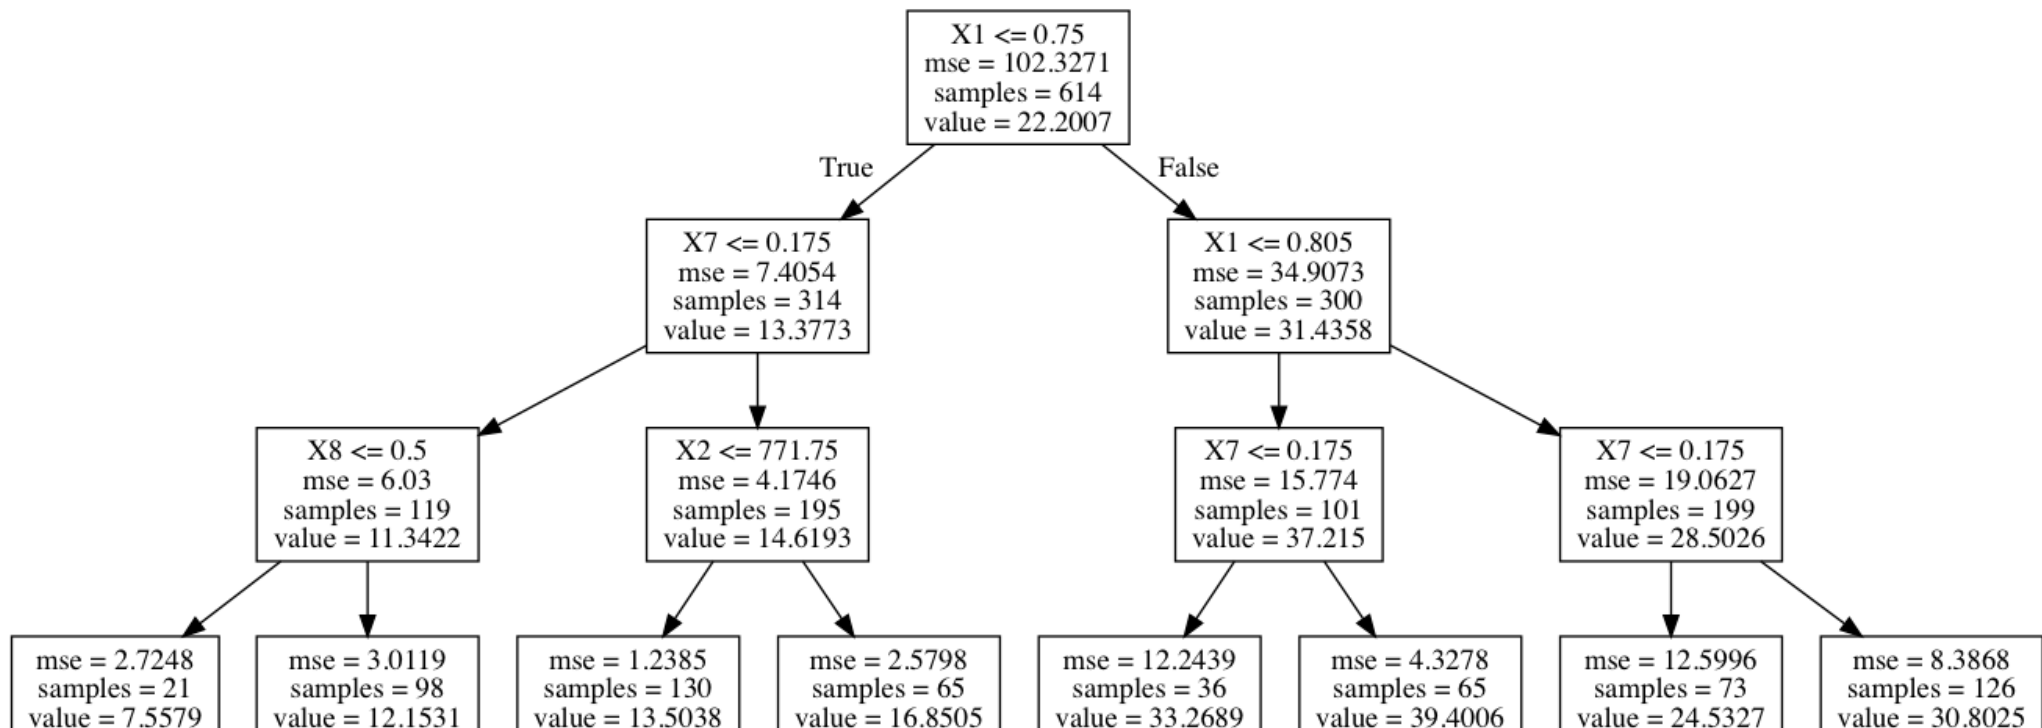
## Big picture concepts:

- Test error can be estimated using bootstrap or approximated with known methods
- Subset selection algorithms help determine a smaller set of  $X$ 's that explain more of the variance in  $Y$
- Regularization and shrinkage methods address the bias/variance tradeoff by adding a penalty that *shrinks* all coefficients toward zero

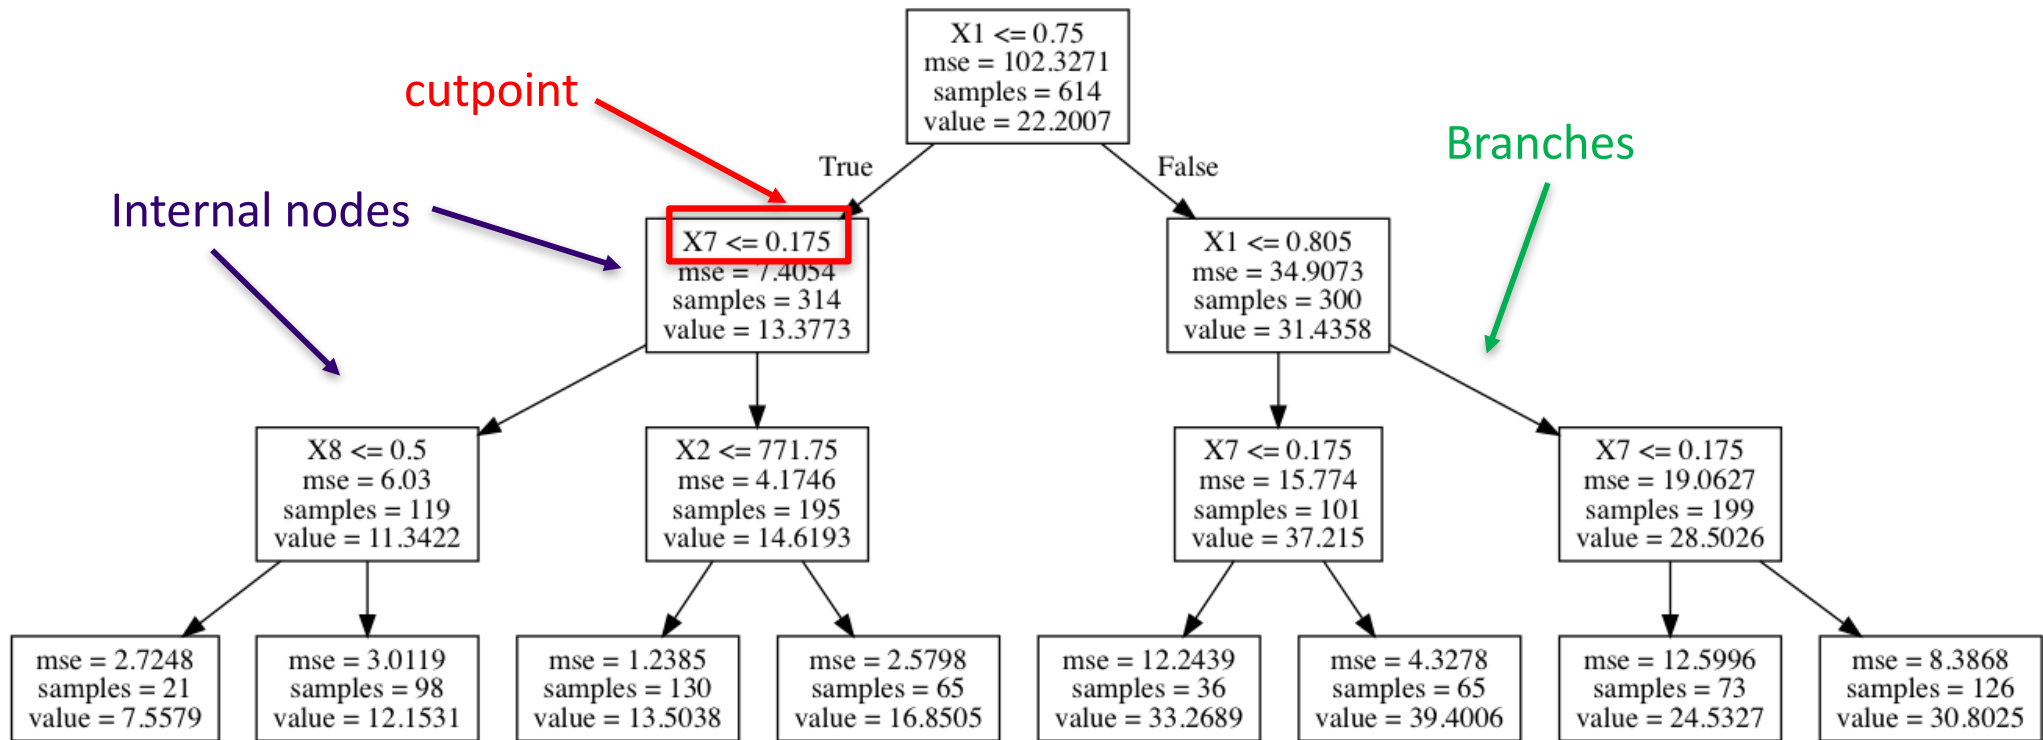


# Decision trees: Example

- > Example of a decision tree that predicts the Y1 predictor from the UCI data set using a simple 3 level regression decision tree



# Decision trees: Nomenclature



Terminal nodes  
or “leaves”

**Nodes:** a point where we make a decision  
(quantitative or qualitative (class) comparisons)

**Branch:** connections between nodes



# Building a regression tree

- > The data in ISL Fig 8.2 shows baseball player salary vs “Hits” vs “Years playing”. The regression tree is divided into three regions (R1-R3) based on sub-divisions of the feature space
- > The goal is to find the cutpoints (s) that minimize the RSS
  - This is done iteratively by first finding the biggest decreases in RSS by appropriately selecting which predictor (j) and cutpoint give the biggest decrease in RSS...

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2, \quad (8.1)$$

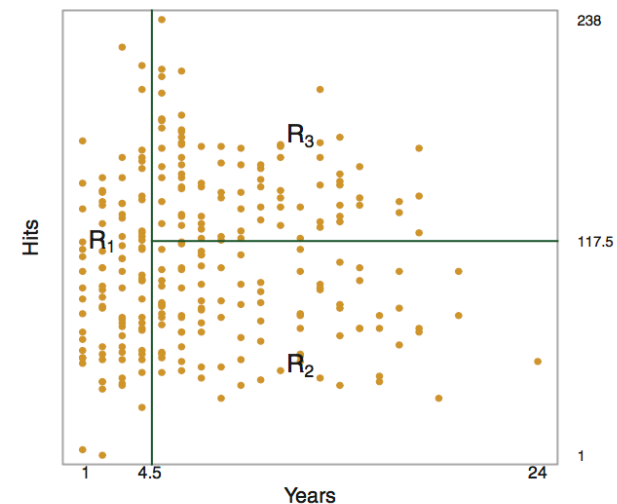
^ RSS for a single region for predictor j

$$R_1(j, s) = \{X | X_j < s\} \text{ and } R_2(j, s) = \{X | X_j \geq s\}, \quad (8.2)$$

^ Cutpoint breaks it into two regions

$$\sum_{i: x_i \in R_1(j, s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i: x_i \in R_2(j, s)} (y_i - \hat{y}_{R_2})^2, \quad (8.3)$$

^ The total RSS is minimized (8.3)



**FIGURE 8.2.** The three-region partition for the *Hitters* data set from the regression tree illustrated in Figure 8.1.

# Options for tree building

---

- > **Approach on the last page is known as “top down” or “greedy” → goes after the biggest reductions in RSS first** (recursive binary splitting)
  - Does not simulate all possible trees, so possible you are not finding the minimum RSS
  - Also possible that trees designed in this approach can lead to overfitting (too much bias)
    - > The procedure of “tree pruning” can address this (not natively supported in sklearn)
- > **Other options/choices in building your DT will be explored in the Python segment...**



# Regression trees vs Classification trees

---

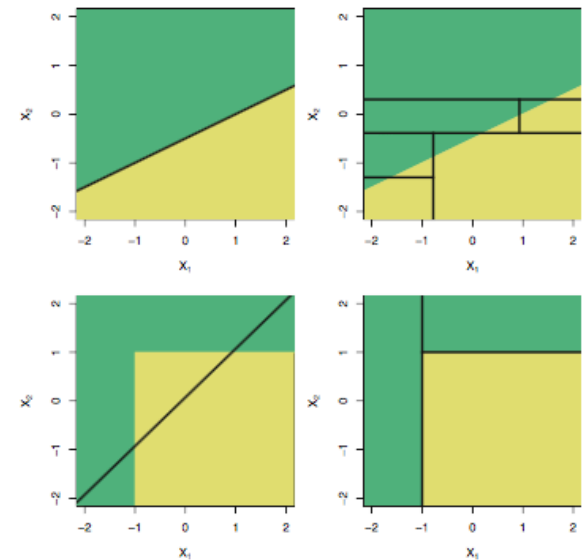
- > Concepts are identical
- > Error metric is based on the classification error rate not RSS , just as in KNN and related methods
- > The DT leaves have your qualitative class assignment, not a quantitative prediction
- > Please see section 8.1.2 of ISL for further information, there are some qualitative differences between the different DT's and assessment of their quality...





# Big picture concepts in building a tree

- > Decision trees are constructed as an alternative to regression models we have seen
- > Fig 8.7 shows some extreme examples of "regression vs. classification" and comparisons
- > Main advantage is ability to handle nonlinear relationships
- > Main disadvantage is high sensitivity to changes in test set



**FIGURE 8.7.** Top Row: A two-dimensional classification example in which the true decision boundary is linear, and is indicated by the shaded regions. A classical approach that assumes a linear boundary (left) will outperform a decision tree that performs splits parallel to the axes (right). Bottom Row: Here the true decision boundary is non-linear. Here a linear model is unable to capture the true decision boundary (left), whereas a decision tree is successful (right).

# In class Python project: Teams of 2

- > You should use the UCI data set and look to predict the Y1 (heating load) response vs all 8 predictors (X1-X8)
  - Break the data into 80% (training), 20% (testing)
- > Use a Python package that can perform decision tree regression
- > Make a plot of testing/training error vs. max tree depth (discuss what you expect the graph to look like first! Really. Draw a sketch!)
  - Don't worry about differentiating validation vs test set data
  - You can hijack a lot of the code from last Wed
- > We will do 2 min standups every 10 mins and go for ~50 mins (I will pick teams or ask for volunteer)
- > If you finish early:
  - Study the properties of your tree at the depth you decided on
    - > Visualize it: ask me for a code snippet
  - Study the sensitivity to the training test set size, random seed, etc.
    - > Test set sensitivity is very important in DT creation!
  - Consider a bootstrap to enable true test MSE estimation
  - Compare error in DT regression to the best fits obtained from MLR, Ridge and LASSO
  - Discuss how you would pick a method to use in the future?



# Building better trees with ensembles

---

- > Big picture ensemble concepts
- > Three common ensemble methods
  - Bagging
  - Random forests
  - **Boosting** [not covered, but similar in qualitative approach]



# Ensemble methods (general)

---

- > **A major takeaway from ISL CH8 is that the error of a DT is highly dependent on the training set used** – sometimes much more than in regression
- > **One way to avoid this is to use so-called ensemble methods**
  - **Conceptually very similar to resampling (bootstrap and cross-validation)**
    - > **Recall main purpose of resampling:** error estimation
  - **In contrast, we go beyond error estimation to use ensemble methods for improved model training**
- > **Important: these methods get introduced in ISL in the context of DTs but can also be applied to other ML techniques**



# Application of ensemble methods to decision trees: Bagging

---

- > The bagging concept builds on bootstrap methods
- > Bagging algorithm
  1. Build a model based on  $B$  individual bootstrap data sets
  2. Make predictions  $f(x)$  for each model and average the predicted response

$$\hat{f}_{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x).$$

- > When applying bagging to the creation of DTs, you should **grow trees with high bias (deep trees)** as the bagging will **reduce** variance
- > **Bagging classifier trees use “majority vote” (most common occurring response)**



# Application of ensemble methods to decision trees: Random Forest

---

- > Bagging methods work well, but the bootstrap sets can still be highly correlated, leading to increased training set error
- > To reduce correlation, the random forest method uses a random sample of predictors at each split (node) in the tree
  - Heuristic is that  $m \approx p^{0.5}$  predictors are randomly chosen at each split
  - The random DT is combined with bagging to create a random forest



# Next week plan

---

- > Unsupervised learning methods, including Python
- > Neural networks OR ML methods for time series data
  - Which do you prefer?
- > Course evaluation and DIRECT update
  - Data Science Option , course numbers , Capstone projects

